



# Introduction to Artificial Intelligence and Machine Learning

School of Engineering  
Nanyang Polytechnic



# Trainer

Foo Yong Wee, Ph.D.

Manager, Computational Intelligence and Analytics Group

School of Engineering

Nanyang Polytechnic

Email: [foo\\_yong\\_wee@nyp.edu.sg](mailto:foo_yong_wee@nyp.edu.sg)


Tel: 6550 0962



# Introduction to Artificial Intelligence and Machine Learning

# Outline

- 1 What is Artificial Intelligence (AI)
- 2 What is Machine Learning (ML)
- 3 What is Deep Learning (DL)
- 4 Applications of AI and ML



# 1 What is AI



# Definition of AI


Artificial Intelligence is the ability of a program or machine to think and behave like a human.

## Turing Test

A machine is said to pass the Turing Test when it exhibits behaviour that is indistinguishable from that of a human.




Alan Turing (1912–1954)



# The Turing Test

A computer passes the test if a human interrogator, after posing some written questions, cannot tell whether the written responses come from a person or from a computer



Artificial Intelligence - The **Turing test**, developed by [Alan Turing](#) in 1950

# History of AI



The 1956 Dartmouth Workshop on Artificial Intelligence  
Workshop on AI at Dartmouth College

"... every aspect of learning or any other feature of intelligence can be so precisely described that machine can be made to simulate it."  
- John McCarthy



Knowledge engineering  
Expert systems  
Computer vision  
Natural language understanding  
Lisp machines  
Japan's fifth generation computer project



Microsoft released the first individual intelligent assistant Microsoft Cortana in the world.



ChatGPT



Watson won Jeopardy game  
Deep Blue beat world chess champion



AlphaGo won Go champion



1943

Artificial neurons (McCullouch and Pitts)

1956

Symbolic processing  
Formal representation (logic, ...)  
Reasoning, inference  
Search and problem solving  
Connectionism (neural networks)

... within a generation the problem of creating artificial intelligence will be substantially solved.

- Marvin Minsky

1974

1980



1987

1993

Data mining  
knowledge discovery  
Machine learning  
Cognitive computing  
Mathematical/statistical methods  
Supercomputers

1997

2011


Big data  
Data analytics  
Massively distributed computing  
Big data distributed databases  
Cloud  
IOT

2016

2020

2023 ...

# AI Resurgence



**BIG DATA, HIGH-PERFORMANCE  
COMPUTING AND CLOUD, HAS  
TRANSFORMED AI**

# AI is All Around Us




Image search



Smart IoT devices



Facial recognition



Recommendation engines



First chess, then Jeopardy, then Go. Now poker too has fallen to AI



Drones, UAV

Self-driving vehicles



Virtual smart assistants

Google  
BERT

ChatGPT



Generative AI



Spam filters



Chatbots



Language translation

Robo-advisors

# AI Technologies



**Computer vision:** a science of how to make computers "see"

**Speech processing:** a general term for various processing technologies used to research the voicing process, statistical features of speech signals, speech recognition, machine-based speech synthesis, and speech perception


**Natural language processing (NLP):** is the use computer technologies to understand and process human language for a range of applications, including machine translation, sentiment analysis, and text analytics.

**Generative AI** generate new, unique data samples that are similar to the examples it was trained on such as images.

**Language Models** a deep learning Generative model that can understand the context and nuances of natural language text.



# Computer Vision



Traffic Analysis

Facial recognition

Comparison Gallery


Authentication result



Electronic Attendance



Common applications include image classification, target detection, image segmentation, target tracking, optical character recognition (OCR), and facial recognition.



Smart Album




Image Search




Authentication



Action Analysis



# Voice/Speech Processing



Question Answering Bot (QABot)



Voice Navigation



Intelligent Education

Other applications:

- Spoken language evaluation
- Diagnostic robot
- Voiceprint recognition
- Smart sound box
- ...




Common applications include voice recognition, voice synthesis, voice wakeup, voiceprint recognition, and audio-based incident detection.



Real-time Conference Records



# Natural Language Processing




## Other applications:

- Knowledge graph
- Intelligent copywriting
- Video subtitle
- ...

Common applications include machine translation, text mining, text summarization, and sentiment analysis.






### Video Generation



Google's Imagen video (<https://Imagen.research.google/video/>) can produce short high-definition video clips using a similar technique to the one it uses to make images. (Photo by Google AI)

# Generative AI




Art Generation


# Generative AI and Large Language Models (LLMs)



# Question

- How can AI technology revolutionize different industries, and what are some specific examples of its applications in these industries?







# End of Chapter 1

## 2 What is Machine Learning

# Can a machine learn?



# Classical Rule-Based Algorithms




- Explicit programming is used to solve problems
- Rules can be manually specified

# Definition of learning




A computer programme is said to **learn** from **experience E** with respect to some class of **tasks T** and **performance measure P**, if its performance at tasks in T, as measured by P, improves with experience E.”

– Tom Michell





# Machine Learning

Machine learning allows computers to learn and infer from data



# Machine Learning Model

After you trained the machine (called a model) to recognize dog or cat, then machine is able to tell your answer when you give it an image



# Machine Learning Vocabulary



The table illustrates the components of machine learning vocabulary:

- Feature:** Sepal length, Sepal width, Petal length, Petal width
- Target:** Species
- Example:** A single row of data.
- Label:** The target value for the example row (Setosa).

Sepal length	Sepal width	Petal length	Petal width	Species
6.7	3.0	5.2	2.3	Virginica
6.4	2.8	5.6	2.1	Virginica
4.6	3.4	1.4	0.3	Setosa
6.9	3.1	4.9	1.5	Versicolor
4.4	2.9	1.4	0.2	Setosa
4.8	3.0	1.4	0.1	Setosa
5.9	3.0	5.1	1.8	Virginica
5.4	3.9	1.3	0.4	Setosa
4.9	3.0	1.4	0.2	Setosa
5.4	3.4	1.7	0.2	Setosa

- **Target:** Predicted category or value of the data (column to predict)
- **Features:** properties of the data used for prediction (non-target columns)
- **Example:** a single data point within the data (one row)
- **Label:** the target value for a single data point (for category), also known as a response value (for numeric)



# Types of Machine Learning

Supervised Learning

Data Points have known outcome


Unsupervised  
Learning

Data Points have unknown outcome

Reinforcement  
Learning

Evaluate actions rather than telling the learning system how to perform correct actions

# Supervised Learning





# Supervised Learning



Duck




Duck



Not Duck



Not Duck



Supervised  
Learning




Predictive  
Model

- The algorithm learns patterns from labelled data and makes predictions and try to label new data


Duck

# Supervised Learning




- Uses labelled training data to improve programs future actions. Reproduces known knowledge. Learn by example approach.
- Supervised learning needs to be given examples of what is “good” and what is “bad”

# UnSupervised Learning




# UnSupervised Learning



- Uses unlabeled data, correct classes are not known. Interpret and Groups the input data only

# Reinforcement Learning

- Close to human learning
- Algorithm (agent) learns a policy of how to act in a given environment
- Every action has some impact in the environment, and the environment provides rewards that guides the learning algorithm
- Reinforcement Learning is learning how to act in order to maximize a numerical reward





# Type of Supervised Learning

Regression

Outcome is continuous (numerical)


Classification

Outcome is a category

Binary


Multiclass

# Supervised Learning Overview






# Supervised Learning



Weather	Temperature	Wind Speed	Enjoy Sports
Sunny	Warm	Strong	Yes
Rainy	Cold	Fair	No
Sunny	Cold	Weak	Yes




# Supervised Learning - Regression

- Regression: To discover the dependencies in the sample dataset by mapping the relationships between the feature values and the response.
- Regression Questions:
  - How much will I benefit from the stock next week?
  - What's the temperature on Tuesday?

Monday	Tuesday
 32°	


# Supervised Learning - Classification

- Classification: maps samples in a sample dataset to a specified category by using a classification model.
- Classification Questions:
  - Will there be a traffic jam on XX road during the morning rush hour tomorrow?
  - Which method is more attractive to customers: \$10 voucher or 25% off?






# Unsupervised Learning



# Examples of Classification Algorithms

- K Nearest Neighbors
- Logistic Regression
- Support Vector Machine
- Gaussian Naive Bayes
- Decision Tree
- Ensemble Methods
  - Random Forest, Gradient Boosting






# Example: Decision Tree

- Decision tree is easy to interpret and implement
- Heterogeneous input data allowed, preprocessing required
- However, decision trees tend to overfit
- Pruning helps reduce variance to a point. Often not significant for model to generalize well


**Temperature >50°F**





## Example: Random Forest


Random Forest classifier is a bagging ensemble method based on lots of decision trees with random selection of subsets of training samples.



The result is based on the majority votes from all the decision trees

# Example: Gradient Boosting

- Gradient Boosting is a ensemble boosting method that "boosting" many weak predictive models into a strong one, in the form of ensemble of weak models.
- Boosting utilizes different loss functions At each stage, the margin is determined for each point. Margin is positive for correctly classified points and negative for misclassifications. Value of loss function is calculated from margin





# Binary Classification

- In many predictive analysis, we are interested in YES/NO analysis such as spam/not-spam, health/not-healthy etc.
- In this case, we can form a **confusion matrix** of 2 columns and 2 rows

	Predicted Positive	Predicted Negative	Type II Error
Actual Positive	True Positive (TP)	False Negative (FN)	
Actual Negative	False Positive (FP)	True Negative (TN)	
			Type I Error

**True Positive:** You predicted positive and it's true

**True Negative:** You predicted negative and it's true

**False Positive:** (Type 1 Error) You predicted positive and it's false

**False Negative:** (Type 2 Error) You predicted negative and it's false.



Several measurements are used to assess the performance of a classifier. These are:  
Accuracy, Precision, Recall, and F1 score.

	Predicted Positive	Predicted Negative
Actual Positive	TP	FN
Actual Negative	FP	TN

# Confusion Matrix

$$\text{Accuracy} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{FN} + \text{FP} + \text{TN}}$$

(is the percentage (x100) of correct hits from all classes – positive and negative)

$$\text{Recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

(is the percentage (x100) of correct hits from all the positive classes)

$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

(is the percentage (x100) of correct hits from all the positive predictions)

$$\text{F1 Score} = \frac{2 \times (\text{Precision} \times \text{Recall})}{\text{Precision} + \text{Recall}}$$

takes both false positive and false negatives into account



# AUC and ROC Curve in Machine Learning

- Receiver Operating Characteristics (ROC) curve provides a graphical representation of the effectiveness of the binary classification model
- It plots the true positive rate (TPR) vs the false positive rate (FPR) at different classification thresholds
- Area Under the Curve (AUC) represents the area under the ROC curve

$$\text{TPR} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

True Positive Rate / Sensitivity / Recall


$$\text{FPR} = \frac{\text{FP}}{\text{TN} + \text{FP}}$$

Also define as  $(1 - \text{Specificity})$



- ROC-AUC tries to measure and ranks the performance of classifiers
- AUC can help you decide which model is better

Sensitivity



# Regression Algorithms

- Linear Regression (Most Common)
- Ridge Regression
- Lasso Regression
- Elastic Net Regression




# What is Linear Regression

- Linear regression is a statistical method of finding the relationship between independent and dependent variables
- In linear regression, we are trying to obtain the line of best fit
- Two common methods to achieve this is:
  - Ordinary Least Square
  - Gradient Descent

# Loss Function and Mean Square Error


- Mean Square Error (MSE) is the common **loss function** to measure how good is the linear regression model
- Machine Learning aims to **minimize** the MSE to find the best linear regression model or the “line of best fit”
- In other words, we want to identify a slope-intercept, in the equation,  $y = mx + b$  (hence the word linear)





# Assessing the Goodness-of-Fit

- After you have fit a linear model using regression analysis, you need to determine how well the model fits the data.
  - R-squared (or Coefficient of Determination)
    - $R^2 = 1 - (\text{Sum of Square Error} / \text{Total Sum of Square})$
    - $0 \leq R^2 \leq 1$
    - If  $R^2 = 0$ , model is not useful, if  $R^2= 1$ , then model fits all points perfectly or the dependent variable(s) is explained by the independent variables
    - A rule of thumb is  $R^2 \geq 0.7$  is a good fit





# Regularization

- Fitting the training data too closely is called overfitting and not fitting it well enough is called underfitting
- Overfitting is a problem in machine learning, and we want to discourage the model from learning overly complex patterns in the data
- Regularization in machine learning is a technique to prevent the model from overfitting by adding extra information, or penalty term, to the model's loss function



# Common Regularization Techniques


- L1 Regularization (Lasso): This technique adds the absolute values of the coefficients as a penalty term. It encourages sparsity in the model, making some features have exactly zero coefficients. It's useful for feature selection.
- L2 Regularization (Ridge): L2 regularization adds the squared values of the coefficients as a penalty term. It doesn't force coefficients to become exactly zero but helps in reducing the magnitude of all coefficients, which can lead to a simpler model.
- Elastic Net Regularization: Elastic Net combines both L1 and L2 regularization. It balances between feature selection (L1) and coefficient shrinkage (L2).

# Clustering Algorithms


- K-Means Clustering
- Hierarchical Agglomerative Clustering




# K-Mean Algorithm




(a)




(b)




(c)



(d)




(e)



(f)


# K-Mean Algorithm Simulation

Click on the images below to start k-mean clustering simulation



<https://www.naftaliharris.com/blog/visualizing-k-means-clustering/>

### 3 What is Deep Learning



# History of Deep Learning Milestones and Tools

## DL Milestones


- 1943: Neural networks
- 1957: Perceptron
- 1974-86: Backpropagation, RBM, RNN
- 1989-98: CNN, MNIST, LSTM, Bidirectional RNN
- 2006: “Deep Learning”, DBN
- 2009: ImageNet
- 2012: AlexNet, Dropout
- 2014: GANs
- 2014: DeepFace
- 2015: ResNet-152
- 2016: AlphaGo
- 2017: AlphaZero, Capsule Networks
- 2018: Google BERT, OpenAI GPT
- 2019: GPT-2
- 2020: GPT-3
- 2021: GPT-3.5

## DL Tools

- Mark 1 Perceptron – 1960
- Torch – 2002
- CUDA – 2007
- Theano – 2008
- Caffe – 2014
- DistBelief – 2011
- TensorFlow 0.1 – 2015
- PyTorch 0.1 – 2017
- TensorFlow 1.0 – 2017
- PyTorch 1.0 – 2017
- TensorFlow 2.0 – 2019
- PyTorch 1.10 – 2021
- TensorFlow 2.7 – 2021

*Full list of acronym can be found in the reference slides*

# The Human Brain and the Neurons




Visualized here are 3% of the neurons and 0.0001% of the synapses in the brain.

Artificial Brain Simulation - Thalamocortical System, 16.7 Million Neurons - 2.1 Billion Synapses  
[https://www.youtube.com/watch?v=PM\\_gTOm9fgk](https://www.youtube.com/watch?v=PM_gTOm9fgk)



# Biological and Artificial Neural Networks



## Human Brain


- 100 billion neurons and 1,000 trillion synapses

## Artificial Neural Network


- ResNet-152: 60 million synapses
- ChatGPT-2: 50 billion neurons
- ChatGPT-3: 60-80 billion neurons and around 100 trillion synapses



# Neuron: Biological Inspiration for Computation



Neuron: computational building block for the brain




(Artificial) Neuron: computational building block for the “neural network”




# Combing Neurons in Hidden Layers: Power to Approximate

Simple Neural Network



● Input Layer

Deep Learning Neural Network




● Hidden Layer


● Output Layer

**Universality:** For any arbitrary function  $f(x)$ , there exists a neural network that closely approximate it for any input  $x$ .

# Deep Learning is Representation Learning (aka feature learning)




# Deep Learning = Learning Hierarchical Representations



Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]





# Convolutional Neural Network Architecture





# CNN Implementation



1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Stride 1


Feature Map

1	0	1
0	1	0
1	0	1

Filter

Convolution filters slides through the image

1x1	1x0	1x1	0	0
0x0	1x1	1x0	1	0
0x1	0x0	1x1	1	1
0	0	1	1	0
0	1	1	0	0

The convolution process

4		
2	4	3
2	3	4

Feature map

4	4
2	3
2	3

Input

2 x 2 Max Pooling


4	4
2	3
2	3

4	4
2	3
2	3

4	4
2	3
2	3



The pooling process

# CNN Implementation (Cont'd)






# Why Deep Learning? Scalable ML




## 4 Applications of AI and ML



# Why not Deep Learning? Real-World Applications



<https://www.youtube.com/watch?v=tF4DML7FIWk>




<https://www.youtube.com/watch?v=868tExoVdQw>





# Image Detection / Localization with Region-Based Methods (Faster R-CNN)

```
ROIs = region_proposal(image)
for ROI in ROIs
    patch = get_patch(image, ROI)
    results = detector(patch)
```




<https://www.youtube.com/watch?v=zrHfvVdvmz4&t=1s>

# Image Detection / Localization with Single Shot Detection Methods (SSD)




# Semantic Segmentation






# Autoencoders



Encoder is a set of convolutional blocks followed by pooling modules that compress the input to the model

The "bottleneck" restricts the flow of information to the decoder from the encoder, allowing only the most vital information to pass through.

Decoder upsampling the convolutional blocks and reconstructs the data back from its encoded form from the bottleneck's output.




# Image Enhancement using AI

- Let's Enhance - <https://letsenhance.io/>
- Topaz Labs - <https://topazlabs.com/>
- Vance AI - <https://vanceai.com/>
- DeepArt.io - <https://deepart.io/>
- Remini - <https://www.remini.ai/>



**Generative Adversarial Networks (GANs)** are a way to make a generative model by having two neural networks compete with each other.




The **discriminator** tries to distinguish genuine data from forgeries created by the generator.

The **generator** turns random noise into imitations of the data, in an attempt to fool the discriminator.



# Generative Adversarial Networks (GANs)



# Generative Adversarial Networks (GANs) (Cont'd)





# Applications of Generative Adversarial Networks (GANs)

Generate Photographs of Human Faces



T. Karras, et al., "Progressive Growing of GANs for Improved Quality, Stability, and Variation", 2017.

Generate Realistic Photographs



A. Brock, et al., "Large Scale GAN Training for High Fidelity Natural Image Synthesis", 2018.

Generate Consistent Video



<https://www.youtube.com/watch?v=9reHvktowLY>

Translation of sketches to color photographs



Phillip Isola, et al., "Image-to-Image Translation with Conditional Adversarial Networks", 2016.



### Generate Cartoon Characters



Y-H Jin, et al., "Towards the Automatic Anime Characters Creation with Generative Adversarial Networks", 2017.

# Applications of Generative Adversarial Networks (GANs)

### Text-to-Image Translation

The small bird has a red head with feathers that fade from red to gray from head to tail



This bird is black with green and has a very short beak



Han Zhang, et al., "StackGAN: Text to Photo-realistic Image Synthesis with Stacked Generative Adversarial Networks", 2016.

# Image Generation using AI

- ArtBreeder - <https://www.artbreeder.com/>
- This Person Does Not Exist -  
<https://thispersondoesnotexist.com/>
- Generated Photos - <https://generated.photos/>
- Deep Dream Generator -  
<https://deeplearningai.org/deep-dream-generator>
- Runway ML - <https://runwayml.com/>

# Music/Lyrics Generation using AI

- <https://www.ampermusic.com/>
- OpenAI Jukebox -  
<https://openai.com/blog/jukebox/>



# Word Embeddings (Word2Vec)




- Word embedding is capturing context of a word in a document, semantic and syntactic similarity, relation with other words
- Word2Vec is a method to construct an embedding. Another way to put it is, they are vector representations of a particular word.
- It is used to predict the source context words (surrounding words) given a target word (the center word).
- E.g. to predict the context [quick, fox] given target word ‘brown’ or [the, brown] given target word ‘quick’

Source Text	Training Samples
The quick brown fox jumps over the lazy dog. →	(the, quick) (the, brown)
The quick brown fox jumps over the lazy dog. →	(quick, the) (quick, brown) (quick, fox)
The quick brown fox jumps over the lazy dog. →	(brown, the) (brown, quick) (brown, fox) (brown, jumps)
The quick brown fox jumps over the lazy dog. →	(fox, quick) (fox, brown) (fox, jumps) (fox, over)

Skip Gram Model



# Recurrent Neural Networks (RNNs)




## Applications

- Sequence Data
- Text
- Speech
- Audio
- Video
- Generation



# Long Term Dependency


Short-term dependence:  
Bob is eating an **apple**




Long-term dependence:  
Bob likes **apples**. He is hungry and decided to have a snack. So now he is eating an **apple**.

**Long Short-Term Memory (LSTM) Networks:**  
Pick What to Forget and What To Remember

# Encoder-Decoder Architecture



## Machine Language Translation






# Language Model – Generative Pre-Train Transformer (GPT)


1. Input embedding
2. Multi-head self-attention
3. Feed-forward network
4. Normalization and residual connections

Ashish Vaswani et. al. 2017. Attention is all you need. In Proceedings of the 31st International Conference on Neural Information Processing Systems (NIPS'17). Curran Associates Inc., Red Hook, NY, USA, 6000–6010.

<https://arxiv.org/abs/1706.03762>



# Deep Reinforcement Learning




Learning to Walk  
via Deep Reinforcement Learning

Submission ID: 60


<https://www.youtube.com/watch?v=n2gE7n11h1Y>



<https://www.facebook.com/watch/?v=344186809644234>



# Towards General AI



Hans Moravec's illustration of the rising tide of the AI capacity increasingly covering the landscape of human competence

A decorative header image featuring a dark blue background with glowing blue lines and nodes, resembling a complex circuit board or neural network.

# Summary

- AI technology is already used in many fields, including medicine, the automotive industry, manufacturing, retail, social media, banking and finance sectors to mention just a few
- Machine learning depends on algorithms or a set of rules or sequence of instructions guiding an operation. The algorithm in ML is learned by searching for patterns in a training dataset that match the labels to the given input values.
- Deep learning uses artificial neural networks inspired by the human brain, to learn from large amounts of data to solve any pattern recognition problem without human intervention.



## Quiz III

- Deep learning is part of a broader family of \_\_\_\_\_ methods based on \_\_\_\_\_ with \_\_\_\_\_ that can be \_\_\_\_\_.
- A. Machine Learning  
B. Supervised or Unsupervised  
C. Representation Learning  
D. Artificial Neural Networks

A decorative background image showing a close-up of a blue printed circuit board (PCB). The PCB features intricate patterns of blue and white conductive traces, capacitors, and other electronic components. The lighting is dramatic, highlighting the metallic textures and creating a futuristic, high-tech feel.

# Acronym References

- CNN Convolutional neural networks
- RBM Restricted Boltzmann machine
- RNN Recurrent neural network
- MNIST Modified National Institute of Standards and Technology
- LSTM Long short-term memory
- DBN Deep belief network
- GAN Generative adversarial network
- BERT Bidirectional Encoder Representations from Transformers
- CUDA Compute Unified Device Architecture
- CNTK Microsoft Cognitive Toolkit
- RCNN Region-based CNN
- HDC Hybrid Dilated Convolution
- DUC Dense Upscaling Convolution



# The End



URL: [Qualtrics Survey](#) | [Qualtrics Experience Management](#)

COURSE RUN ID: 793880



Course Title:

**INTRODUCTION TO AI & MACHINE LEARNING**



Course Start Date: 5 Oct 2023



Name of Trainer(s): Dr Foo Yong Wee



URL: <https://forms.gle/NBo6VvmeTrqqxvKV8>

Course Title: **INTRODUCTION TO AI & MACHINE LEARNING  
(CE1574-230001)**

Course Start Date: 5 Oct 2023

Name of Trainer(s): Dr Foo Yong Wee